

## Data Preparation

```
ab_path = "/kaggle/input/image-colorization/ab/ab/ab1.npy"
l_path = "/kaggle/input/image-colorization/l/gray_scale.npy"
ab_df = np.load(ab_path)[0:5000]
L_df = np.load(l_path)[0:5000]
dataset = (L_df, ab_df)
gc.collect()
```

## Data Exploration and Visualiaztion

```
def lab_to_rgb(L, ab):
    """
    Takes an image or a batch of images and converts from LAB space to RGB
    """
    L = L * 100
    ab = (ab - 0.5) * 128 * 2
    Lab = torch.cat([L, ab], dim=2).numpy()
    rgb_imgs = []
    for img in Lab:
        img_rgb = lab2rgb(img)
        rgb_imgs.append(img_rgb)
    return np.stack(rgb_imgs, axis=0)

plt.figure(figsize=(30,30))
for i in range(1,16,2):
    plt.subplot(4,4,i)
    img = np.zeros((224,224,3))
    img[:, :, 0] = L_df[i]
    plt.title('B&W')
    plt.imshow(lab2rgb(img))

    plt.subplot(4,4,i+1)
```

```

img[:, :, 1:] = ab_df[i]
img = img.astype('uint8')
img = cv2.cvtColor(img, cv2.COLOR_LAB2RGB)
plt.title('Colored')
plt.imshow(img)
gc.collect()

```

## Data Loader

```

class ImageColorizationDataset(Dataset):
    """ Black and White (L) Images and corresponding A&B Colors"""
    def __init__(self, dataset, transform=None):
        """
        :param dataset: Dataset name.
        :param data_dir: Directory with all the images.
        :param transform: Optional transform to be applied on sample
        """
        self.dataset = dataset
        self.transform = transform
    def __len__(self):
        return len(self.dataset[0])
    def __getitem__(self, idx):
        L = np.array(dataset[0][idx]).reshape((224,224,1))
        L = transforms.ToTensor()(L)
        ab = np.array(dataset[1][idx])
        ab = transforms.ToTensor()(ab)
        return ab, L

batch_size = 1

# Prepare the Datasets
train_dataset = ImageColorizationDataset(dataset = (L_df, ab_df))
test_dataset = ImageColorizationDataset(dataset = (L_df, ab_df))

```

```
# Build DataLoaders
```

```
train_loader = DataLoader(dataset=train_dataset, batch_size=batch_size, shuffle = True,  
pin_memory = True)
```

```
test_loader = DataLoader(dataset=test_dataset, batch_size=batch_size, shuffle = False,  
pin_memory = True)
```

## Data Modelling

```
class ResBlock(nn.Module):
```

```
    def __init__(self, in_channels, out_channels, stride=1):
```

```
        super().__init__()
```

```
        self.layer = nn.Sequential(
```

```
            nn.Conv2d(in_channels,out_channels,kernel_size=3, padding=1, stride=stride,  
bias=False),
```

```
            nn.BatchNorm2d(out_channels),
```

```
            nn.ReLU(inplace=True),
```

```
            nn.Conv2d(out_channels, out_channels,kernel_size=3,padding=1, stride=1, bias=False),
```

```
            nn.BatchNorm2d(out_channels),
```

```
            nn.ReLU(inplace=True)
```

```
        )
```

```
        self.identity_map = nn.Conv2d(in_channels, out_channels,kernel_size=1,stride=stride)
```

```
        self.relu = nn.ReLU(inplace=True)
```

```
    def forward(self, inputs):
```

```
        x = inputs.clone().detach()
```

```
        out = self.layer(x)
```

```
        residual = self.identity_map(inputs)
```

```
        skip = out + residual
```

```
        return self.relu(skip)
```

```
class DownSampleConv(nn.Module):
```

```
    def __init__(self, in_channels, out_channels, stride=1):
```

```

        super().__init__()
        self.layer = nn.Sequential(
            nn.MaxPool2d(2),
            ResBlock(in_channels, out_channels)
        )
    def forward(self, inputs):
        return self.layer(inputs)

class UpSampleConv(nn.Module):
    def __init__(self, in_channels, out_channels):
        super().__init__()
        self.upsample = nn.Upsample(scale_factor=2, mode="bilinear", align_corners=True)
        self.res_block = ResBlock(in_channels + out_channels, out_channels)
    def forward(self, inputs, skip):
        x = self.upsample(inputs)
        x = torch.cat([x, skip], dim=1)
        x = self.res_block(x)
        return x

class Generator(nn.Module):
    def __init__(self, input_channel, output_channel, dropout_rate = 0.2):
        super().__init__()
        self.encoding_layer1_ = ResBlock(input_channel, 64)
        self.encoding_layer2_ = DownSampleConv(64, 128)
        self.encoding_layer3_ = DownSampleConv(128, 256)
        self.bridge = DownSampleConv(256, 512)
        self.decoding_layer3_ = UpSampleConv(512, 256)
        self.decoding_layer2_ = UpSampleConv(256, 128)
        self.decoding_layer1_ = UpSampleConv(128, 64)
        self.output = nn.Conv2d(64, output_channel, kernel_size=1)
        self.dropout = nn.Dropout2d(dropout_rate)

```

```

def forward(self, inputs):
    ##### Encoder #####
    e1 = self.encoding_layer1_(inputs)
    e1 = self.dropout(e1)
    e2 = self.encoding_layer2_(e1)
    e2 = self.dropout(e2)
    e3 = self.encoding_layer3_(e2)
    e3 = self.dropout(e3)

    ##### Bridge #####
    bridge = self.bridge(e3)
    bridge = self.dropout(bridge)

    ##### Decoder #####
    d3 = self.decoding_layer3_(bridge, e3)
    d2 = self.decoding_layer2_(d3, e2)
    d1 = self.decoding_layer1_(d2, e1)

    ##### Output #####
    output = self.output(d1)

    return output

model = Generator(1,2).to(device)
summary(model, (1, 224, 224), batch_size = 1)

```

## B. Discriminator ( Critic )

```

class Critic(nn.Module):
    def __init__(self, in_channels=3):
        super(Critic, self).__init__()

        def critic_block(in_filters, out_filters, normalization=True):
            """Returns layers of each critic block"""

            layers = [nn.Conv2d(in_filters, out_filters, 4, stride=2, padding=1)]

            if normalization:

```

```

        layers.append(nn.InstanceNorm2d(out_filters))
    layers.append(nn.LeakyReLU(0.2, inplace=True))
    return layers

self.model = nn.Sequential(
    *critic_block(in_channels, 64, normalization=False),
    *critic_block(64, 128),
    *critic_block(128, 256),
    *critic_block(256, 512),
    nn.AdaptiveAvgPool2d(1),
    nn.Flatten(),
    nn.Linear(512, 1)
)

def forward(self, ab, l):
    # Concatenate image and condition image by channels to produce input
    img_input = torch.cat((ab, l), 1)
    output = self.model(img_input)
    return output

model = Critic(3).to(device)
summary(model, [(2, 224, 224), (1, 224, 224)], batch_size = 1)

```

### C. Generative Adversarial Network

```

def _weights_init(m):
    if isinstance(m, (nn.Conv2d, nn.ConvTranspose2d)):
        torch.nn.init.normal_(m.weight, 0.0, 0.02)
    if isinstance(m, nn.BatchNorm2d):
        torch.nn.init.normal_(m.weight, 0.0, 0.02)
        torch.nn.init.constant_(m.bias, 0)

def display_progress(cond, real, fake, current_epoch = 0, figsize=(20,15)):
    """
    Save cond, real (original) and generated (fake)
    """

```

images in one panel

"""

```
cond = cond.detach().cpu().permute(1, 2, 0)
```

```
real = real.detach().cpu().permute(1, 2, 0)
```

```
fake = fake.detach().cpu().permute(1, 2, 0)
```

```
images = [cond, real, fake]
```

```
titles = ['input','real','generated']
```

```
print(f'Epoch: {current_epoch}')
```

```
fig, ax = plt.subplots(1, 3, figsize=figsize)
```

```
for idx,img in enumerate(images):
```

```
    if idx == 0:
```

```
        ab = torch.zeros((224,224,2))
```

```
        img = torch.cat([images[0]* 100, ab], dim=2).numpy()
```

```
        imgan = lab2rgb(img)
```

```
    else:
```

```
        imgan = lab_to_rgb(images[0],img)
```

```
    ax[idx].imshow(imgan)
```

```
    ax[idx].axis("off")
```

```
for idx, title in enumerate(titles):
```

```
    ax[idx].set_title('{}'format(title))
```

```
plt.show()
```

```
class CWGAN(pl.LightningModule):
```

```
    def __init__(self, in_channels, out_channels, learning_rate=0.0002, lambda_recon=100,
display_step=10, lambda_gp=10, lambda_r1=10,):
```

```
        super().__init__()
```

```
        self.save_hyperparameters()
```

```
        self.display_step = display_step
```

```
        self.generator = Generator(in_channels, out_channels)
```

```
        self.critic = Critic(in_channels + out_channels)
```

```

self.optimizer_G = optim.Adam(self.generator.parameters(), lr=learning_rate, betas=(0.5,
0.9))

self.optimizer_C = optim.Adam(self.critic.parameters(), lr=learning_rate, betas=(0.5, 0.9))

self.lambda_recon = lambda_recon

self.lambda_gp = lambda_gp

self.lambda_r1 = lambda_r1

self.recon_criterion = nn.L1Loss()

self.generator_losses, self.critic_losses = [], []

def configure_optimizers(self):

    return [self.optimizer_C, self.optimizer_G]

def generator_step(self, real_images, conditioned_images):

    # WGAN has only a reconstruction loss

    self.optimizer_G.zero_grad()

    fake_images = self.generator(conditioned_images)

    recon_loss = self.recon_criterion(fake_images, real_images)

    recon_loss.backward()

    self.optimizer_G.step()

    # Keep track of the average generator loss

    self.generator_losses += [recon_loss.item()]

def critic_step(self, real_images, conditioned_images):

    self.optimizer_C.zero_grad()

    fake_images = self.generator(conditioned_images)

    fake_logits = self.critic(fake_images, conditioned_images)

    real_logits = self.critic(real_images, conditioned_images)

    # Compute the loss for the critic

    loss_C = real_logits.mean() - fake_logits.mean()

    # Compute the gradient penalty

    alpha = torch.rand(real_images.size(0), 1, 1, 1, requires_grad=True)

    alpha = alpha.to(device)

    interpolated = (alpha * real_images + (1 - alpha) * fake_images.detach()).requires_grad_(True)

```



```

interpolated_logits = self.critic(interpolated, conditioned_images)

grad_outputs = torch.ones_like(interpolated_logits, dtype=torch.float32, requires_grad=True)

gradients = torch.autograd.grad(outputs=interpolated_logits, inputs=interpolated,
grad_outputs=grad_outputs,create_graph=True, retain_graph=True)[0]

gradients = gradients.view(len(gradients), -1)

gradients_penalty = ((gradients.norm(2, dim=1) - 1) ** 2).mean()

loss_C += self.lambda_gp * gradients_penalty

# Compute the R1 regularization loss

r1_reg = gradients.pow(2).sum(1).mean()

loss_C += self.lambda_r1 * r1_reg

# Backpropagation

loss_C.backward()

self.optimizer_C.step()

self.critic_losses += [loss_C.item()]

def training_step(self, batch, batch_idx, optimizer_idx):

    real, condition = batch

    if optimizer_idx == 0:

        self.critic_step(real, condition)

    elif optimizer_idx == 1:

        self.generator_step(real, condition)

    gen_mean = sum(self.generator_losses[-self.display_step:]) / self.display_step

    crit_mean = sum(self.critic_losses[-self.display_step:]) / self.display_step

    if self.current_epoch%self.display_step==0 and batch_idx==0 and optimizer_idx==1:

        fake = self.generator(condition).detach()

        torch.save(cwgan.generator.state_dict(), "ResUnet_"+ str(self.current_epoch) + ".pt")

        torch.save(cwgan.critic.state_dict(), "PatchGAN_"+ str(self.current_epoch) + ".pt")

        print(f"Epoch {self.current_epoch}: Generator loss: {gen_mean}, Critic loss: {crit_mean}")

        display_progress(condition[0], real[0], fake[0], self.current_epoch)

```

```

gc.collect()

cwgan = CWGAN(in_channels = 1, out_channels = 2, learning_rate=2e-4, lambda_recon=100,
display_step=10)

trainer = pl.Trainer(max_epochs=150, gpus=-1)

trainer.fit(cwgan, train_loader)

plt.figure(figsize=(30,60))

idx = 1

for batch_idx, batch in enumerate(test_loader):

    real, condition = batch

    pred = cwgan.generator(condition).detach().squeeze().permute(1, 2, 0)

    condition = condition.detach().squeeze(0).permute(1, 2, 0)

    real = real.detach().squeeze(0).permute(1, 2, 0)

    plt.subplots_adjust(wspace=0, hspace=0)

    plt.subplot(6,3,idx)

    plt.grid(False)

    ab = torch.zeros((224,224,2))

    img = torch.cat([condition * 100, ab], dim=2).numpy()

    imgan = lab2rgb(img)

    plt.imshow(imgan)

    plt.title('Input')

    plt.subplot(6,3,idx + 1)

    ab = torch.zeros((224,224,2))

    imgan = lab_to_rgb(condition,real)

    plt.imshow(imgan)

    plt.title('Real')

    plt.subplot(6,3,idx + 2)

    imgan = lab_to_rgb(condition,pred)

    plt.title('Generated')

    plt.imshow(imgan)

    idx += 3

```

```

if idx >= 18:
    break
torch.set_grad_enabled(False)
cwgan.generator.eval()
all_preds = []
all_real = []
for batch_idx, batch in enumerate(test_loader):
    real, condition = batch
    pred = cwgan.generator(condition).detach()
    Lab = torch.cat([condition, pred], dim=1).numpy()
    Lab_real = torch.cat([condition, real], dim=1).numpy()
    all_preds.append(Lab.squeeze())
    all_real.append(Lab_real.squeeze())
    if batch_idx == 500:
        break
class InceptionScore:
    def __init__(self, device):
        self.device = device
        self.inception = inception_v3(pretrained=True, transform_input=False).to(self.device)
        self.inception.eval()
    def calculate_is(self, generated_images):
        generated_images = generated_images.to(self.device)
        with torch.no_grad():
            generated_features = self.inception(generated_images.view(-1,3,224,224))
            generated_features = generated_features.view(generated_features.size(0), -1)
            p = F.softmax(generated_features, dim=1)
            kl = p * (torch.log(p) - torch.log(torch.tensor(1.0/generated_features.size(1)).to(self.device)))
            kl = kl.sum(dim=1)

```

```

        return kl.mean().item(), kl.std().item()
device = "cuda" # or "cpu" if you don't have a GPU
is_calculator = InceptionScore(device)
all_preds = np.concatenate(all_preds, axis=0)
all_preds = torch.tensor(all_preds).float()
all_real = np.concatenate(all_real, axis=0)
all_real = torch.tensor(all_real).float()
is_model = InceptionScore(device)
# Calculate the Inception Score
mean_real, std_real = is_model.calculate_is(all_real)
print("Inception Score of real images: mean: {:.4f}, std: {:.4f}".format(mean_real, std_real))
mean_is, std_is = is_model.calculate_is(all_preds)
print("Inception Score of fake images: mean: {:.4f}, std: {:.4f}".format(mean_is, std_is))

```

## B. Implemenation of Fréchet Inception Distance (FID)

```

class FID:
    def __init__(self, device):
        self.device = device

        self.inception = inception_v3(pretrained=True, transform_input=False).to(self.device)
        self.inception.eval()

        self.mu = None
        self.sigma = None

    def calculate_fid(self, real_images, generated_images):
        real_images = real_images.to(self.device)
        generated_images = generated_images.to(self.device)

        with torch.no_grad():
            real_features = self.inception(real_images.view(-1,3,224,224))
            generated_features = self.inception(generated_images.view(-1,3,224,224))

```

```

real_features = real_features.view(real_features.size(0), -1)
generated_features = generated_features.view(generated_features.size(0), -1)

if self.mu is None:
    self.mu = real_features.mean(dim=0)

if self.sigma is None:
    self.sigma = real_features.std(dim=0)

real_mu = real_features.mean(dim=0)
real_sigma = real_features.std(dim=0)

generated_mu = generated_features.mean(dim=0)
generated_sigma = generated_features.std(dim=0)

mu_diff = real_mu - generated_mu
sigma_diff = real_sigma - generated_sigma

fid = mu_diff.pow(2).sum() + (self.sigma - generated_sigma).pow(2).sum() + (self.mu -
generated_mu).pow(2).sum()

return fid.item()

# Initialize the FID class

device = "cuda" # or "cpu" if you don't have a GPU

fid_calculator = FID(device)

# Calculate the FID

fid_value = fid_calculator.calculate_fid(all_real, all_preds)

print("FID: {:.4f}".format(fid_value))

```

